

경량 로그 기반 지연 업데이트 기법을 활용한 리눅스 커널 확장성 향상

A Lightweight Log-based Deferred Update for
Linux Kernel Scalability

Joohyun Kyong


School of Computer Science

Kookmin University

Thesis advisor: Sung-soo Lim


November 8, 2016

40 Years of Microprocessor Trend Data



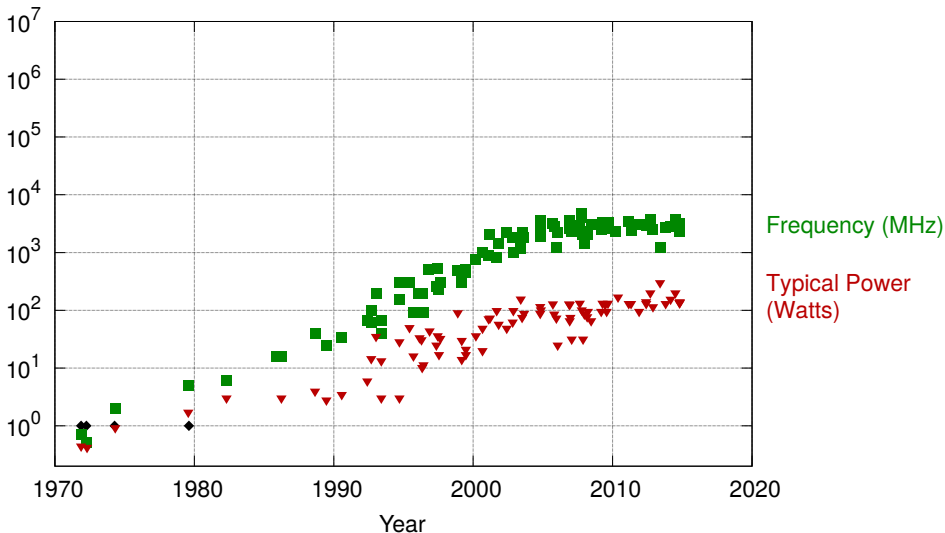
Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2015 by K. Rupp

40 Years of Microprocessor Trend Data




Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2015 by K. Rupp

40 Years of Microprocessor Trend Data




Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2015 by K. Rupp

40 Years of Microprocessor Trend Data




Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2015 by K. Rupp

40 Years of Microprocessor Trend Data




Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2015 by K. Rupp


Performance Scalability



Performance Scalability



Performance Scalability



OS Kernel Scalability

- ▶ OS kernel scalability is an important part for the whole the system parallelism.
- ▶ If the kernel does not scale, applications will not scale.

OS Kernel Scalability

- ▶ OS kernel scalability is an important part for the whole the system parallelism.
- ▶ If the kernel does not scale, applications will not scale.

AIM7 Scalability

Lock Contention Problem

Update Serialization

OS Kernel Scalability History

OS Kernel Scalability History

OS Kernel Scalability History

OS Kernel Scalability History

OS Kernel Scalability History

High update rate Data Structuer

Update-heavy Data structure

Update-heavy Data structure

Non-blocking Data Structure

Non-blocking Data Structure

Non-blocking Data Structure

Non-blocking Data Structure

Cache communication bottlenect

Cache communication bottlenect

Cache communication bottlenect

Cache communication bottlenect

Cache communication bottlenect

Approach: Log-based Concurrent Update

Approach: Log-based Concurrent Update

Approach: Log-based Concurrent Update

Advantages of eliminating time-stamp counters

Contributions

- ▶ Background of research
- ▶ Our new method and Evaluation
- ▶ Future plans and Summary

Outline

- ▶ Design
 - ▶ Approach
 - ▶ Example
- ▶ Applying the Linux kernel
- ▶ Implementation
- ▶ Evaluation

Design

Why the OpLog needs the time-stamp counter?

Log Example

Update-side removing

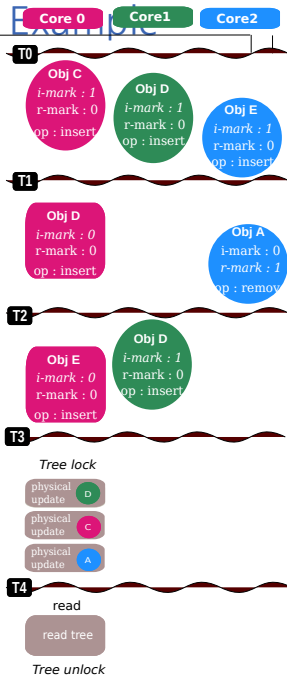
Garbage log

Reusing logs

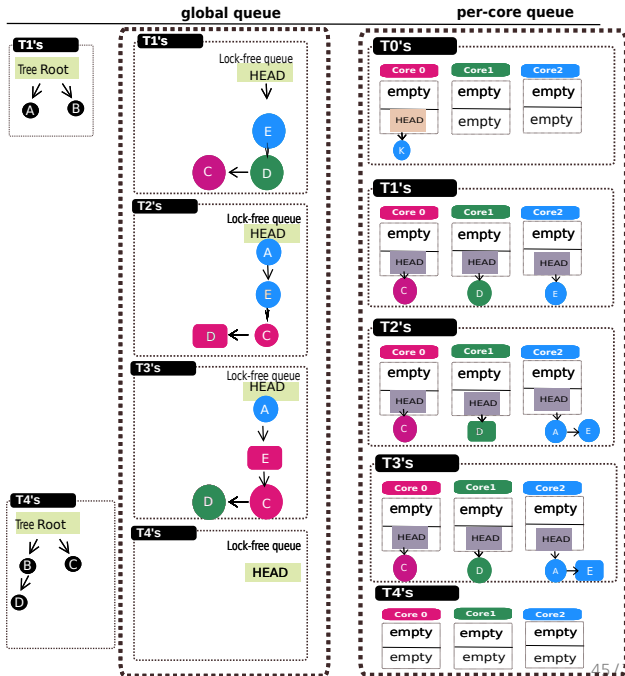
Approach

Example

Threads



memory



Example

Example

Example

Example

Example

Example

Example

Applying the Linux kernel

anonymous reverse mapping

anonymous reverse mapping

file mapping

file mapping

Evaluation

Non-blocking algorithm – Harris linked list

Test-bed

Test-bed

AIM7 – CPU utilization

EXIM – CPU utilization

Lmbench

Lmbench – CPU utilization

Update ratio

Related work

Papers

Conclusion

- ▶ Background of research
- ▶ LDU method and Evaluation
- ▶ Future plans and Summary

Conclusion

- ▶ Background of research
- ▶ LDU method and Evaluation
- ▶ Future plans and Summary
- ▶ <https://github.com/KMU-embedded/scalablelinux>